

GPU Accelerated Poisson solver

A comparison of PyTorch, Cupy and multi-GPU
(Cupy) for solving the Poisson equation

Sebastian Miles, Bingzhou Xie

Supervisor:
Lars Davidson



CHALMERS

CHALMERS UNIVERSITY OF TECHNOLOGY
FEBRUARY 16, 2025

1 Abstract

The purpose of this project is to accelerate a CPU Poisson equation solver with a GPU. The basis of our code is the vectorized Poisson Solver by Lars Davidson [1]. This report documents the steps taken to convert the existing CPU Poisson solver to run efficiently on a GPU. We try two different libraries for this: PyTorch and Cupy and compare their differences and conclude the one that works the best for this particular problem. Finally we also try comparing with multiple GPUs for a single problem. Comparing CPU and PyTorch, there was roughly a 6x speedup with PyTorch. Between PyTorch and Cupy there was a better time complexity with Cupy, where Cupy had support for efficient CSR format. In our test with a 1000×1000 grid, Cupy was running 9x faster than PyTorch.

1.1 Contributions

- Sebastian Miles: Sections 2-5 & 7
- Bingzhou Xie: Sections 6 & 8

2 Solver

When solving the Poisson equation, our main interest is finding an appropriate method for solving the linear matrix equation

$$\mathbf{Ax} = \mathbf{b} \quad (1)$$

where A is a sparse, symmetric and positive-definite matrix.

From the initial code by lada, we implemented our own solvers. We started with the Conjugate Gradient (CG) method. After noticing a very slow convergence rate for large grid sizes, we decided to try preconditioned CG. In particular a Jacobi preconditioner, which is done by extracting only the diagonal of the matrix and leaving everything else to zero. With a tolerance of 10^{-10} on a 100×100 grid the original CG converged after 1173 iterations, while the preconditioned CG converged after 18 iterations.

2.1 Conjugate Gradient (CG)

Suppose that

$$P = \{\mathbf{p}_1, \dots, \mathbf{p}_n\}$$

is a set of n vectors such that $\mathbf{p}_i^T \mathbf{A} \mathbf{p}_j = 0$ for all $i \neq j$, where \mathbf{A} is the same as in equation 1. Then P forms a basis on \mathbb{R}^n . This means that we can express a

solution $\mathbf{x} = \hat{\mathbf{x}}$ of $\mathbf{Ax} = \mathbf{b}$ as

$$\hat{\mathbf{x}} = \sum_i^n c_i \mathbf{p}_i.$$

Left multiplying by A and then by \mathbf{p}_k^T yields

$$\mathbf{A}\hat{\mathbf{x}} = \sum_i^n c_i \mathbf{Ap}_i \implies \mathbf{p}_k^T \mathbf{A}\hat{\mathbf{x}} = \sum_i^n c_i \mathbf{p}_k^T \mathbf{Ap}_i = c_k \mathbf{p}_k^T \mathbf{Ap}_k.$$

This means that

$$c_k = \frac{\mathbf{p}_k^T \mathbf{A}\hat{\mathbf{x}}}{\mathbf{p}_k^T \mathbf{Ap}_k}.$$

The last equality holds because $\mathbf{p}_i^T \mathbf{Ap}_j = 0$ for all $i \neq j$ as part of our premise. This gives us a brief idea for solving equation 1: Find the set of vectors P , then compute the coefficients c_k and then rebuild $\hat{\mathbf{x}}$ from the linear combination.

Consider the quadratic function

$$f : \mathbb{R}^n \rightarrow \mathbb{R}$$

$$f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T \mathbf{Ax} - \mathbf{x}^T \mathbf{b}.$$

We note that the gradient is equal to

$$\nabla f(\mathbf{x}) = \mathbf{Ax} - \mathbf{b}$$

and that $\hat{\mathbf{x}}$ is in fact a minimizer to this function, since that the hessian matrix is positive definite

$$\mathbf{H}(f(x)) = \mathbf{A}.$$

By the Gradient descent method, f decreases the fastest in the direction of the negative gradient $-\nabla f(\mathbf{x})$. Thus, starting with an initial value of \mathbf{x}_0 we set the first vector to the negative gradient $\mathbf{p}_0 = \mathbf{b} - \mathbf{Ax}_0$. Let \mathbf{r}_k be the residue at the k -th step:

$$\mathbf{r}_k = \mathbf{b} - \mathbf{Ax}_k$$

To enforce the conjugation constraint on \mathbf{p}_k we set it to

$$\mathbf{p}_k = \mathbf{r}_k - \sum_{i < k} \frac{\mathbf{p}_i^T \mathbf{Ar}_k}{\mathbf{p}_i^T \mathbf{Ap}_i} \mathbf{p}_i.$$

Continuing along, the next iterations are

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$$

where α_k is some scalar. Let $g(\alpha_k) = f(\mathbf{x}_k + \alpha_k \mathbf{p}_k)$. We can construct a heuristically good scalar with $g'(\alpha_k) \rightarrow 0$ as $k \rightarrow \infty$.

$$\begin{aligned}
g(\alpha_k) &= \frac{1}{2}(\mathbf{x}_k + \alpha_k \mathbf{p}_k)^T \mathbf{A}(\mathbf{x}_k + \alpha_k \mathbf{p}_k) - (\mathbf{x}_k + \alpha_k \mathbf{p}_k)^T \mathbf{b} \\
\implies g'(\alpha_k) &= \frac{1}{2} \mathbf{x}_k^T \mathbf{A} \mathbf{p}_k + \frac{1}{2} \mathbf{p}_k^T \mathbf{A} \mathbf{x}_k + \alpha_k \mathbf{p}_k^T \mathbf{A} \mathbf{p}_k - \mathbf{p}_k^T \mathbf{b} \\
g'(\alpha_k) \rightarrow 0 &\iff \alpha_k \mathbf{p}_k^T \mathbf{A} \mathbf{p}_k \rightarrow \mathbf{p}_k^T (b - \frac{1}{2} \mathbf{A} \mathbf{x}_k) - \frac{1}{2} \mathbf{x}_k^T \mathbf{A} \mathbf{p}_k \\
&= \mathbf{p}_k^T (b - \frac{1}{2} \mathbf{A} \mathbf{x}_k) - \frac{1}{2} (\mathbf{x}_{k+1} - \alpha_k \mathbf{p}_k)^T \mathbf{A} \left(\frac{\mathbf{x}_{k+1} - \mathbf{x}_k}{\alpha_k} \right) \\
&= \mathbf{p}_k^T (b - \mathbf{A} \mathbf{x}_k) - \frac{\mathbf{x}_{k+1}^T \mathbf{A}}{2\alpha_k} (x_{k+1} - x_k) + \frac{\mathbf{p}_k^T \mathbf{A}}{2} (\mathbf{x}_{k+1} - \mathbf{x}_k) \\
&\approx \mathbf{p}_k^T (b - \mathbf{A} \mathbf{x}_k) = \mathbf{p}_k^T r_k
\end{aligned}$$

After dividing by $\mathbf{p}_k^T \mathbf{A} \mathbf{p}_k$ on both sides, we end up with the scalar

$$\alpha_k = \frac{\mathbf{p}_k^T \mathbf{r}_k}{\mathbf{p}_k^T \mathbf{A} \mathbf{p}_k}.$$

2.2 Jacobi Preconditioned Conjugate Gradient

A preconditioning \mathbf{M} means that we transform the matrix into $\mathbf{M}^{-1} \mathbf{A}$ which is more suitable for numerical calculations. A metric commonly used for this is the condition number $\kappa(\mathbf{A})$. The metric measures how much the output of $\mathbf{A}\mathbf{x}$ can change for a small change in \mathbf{x} . The goal of preconditioning is reducing the condition number, that is $\kappa(\mathbf{M}^{-1} \mathbf{A}) < \kappa(\mathbf{A})$. The Jacobi preconditioner is one of the more simpler preconditions. It is defined as $\mathbf{M} = \text{diag}(\mathbf{A})$. Note that the inverse \mathbf{M}^{-1} is a diagonal matrix with the diagonal $M_{i,i}^{-1} = \frac{1}{M_{i,i}}$. So the transformation is computationally cheap.

3 Profiling

The line_profiler package in python is a useful tool for profiling the code line by line. We started by profiling on the CPU, and came to the conclusion that the main time sink was within the CG method. In figure 2 the heaviest line is #200 which takes up 90.8% of the time. The function `A_mv(p)` is a matrix vector product and it would be a good idea to parallelize this operation.

Line #	Hits	Time	Per Hit	% Time	Line Contents
187					@profile
188					def conjugate_gradient(A_mv, b, x0, tol=1e-6, max_iter=1000):
189	1	10.0	10.0	0.0	x = x0
190	1	14320.0	14320.0	0.7	r = b - A_mv(x)
191					# Check if we can exit out early
192					if r @ r < tol:
193	1	404.0	404.0	0.0	return x
194					195
196	1	8.0	8.0	0.0	p = r
197	1	160.0	160.0	0.0	r_norm_sq = r @ r
198					199
199	101	1477.0	14.6	0.1	for _ in range(max_iter):
200	100	1813947.0	18139.5	90.8	Ap = A_mv(p)
201	100	35667.0	356.7	1.8	alpha = r_norm_sq / (p @ Ap)
202	100	28772.0	287.7	1.4	x = x + alpha * p
203	100	25046.0	250.5	1.3	r_new = r - alpha * Ap
204	100	10582.0	105.8	0.5	r_norm_sq_new = r_new @ r_new # dot product of error
205	100	11954.0	119.5	0.6	beta = r_norm_sq_new / r_norm_sq
206	100	2570.0	25.7	0.1	r = r_new
207	100	2298.0	23.0	0.1	r_norm_sq = r_norm_sq_new
208					209
209	100	24420.0	244.2	1.2	if r_norm_sq < tol:
210					break
211					212
212	100	26502.0	265.0	1.3	p = r + beta * p
213					214
214	1	4.0	4.0	0.0	return x
Total time: 0.0721684 s					

Figure 1: Line profiler on the CG function

4 PyTorch Implementation

From the CPU poisson solver by lada, we converted all of the numpy code to PyTorch. With this implementation we could easily swap between a CPU or CUDA device with a single boolean variable. We decided to put all of the PyTorch tensors on the CUDA device as there would be no time spent moving data between the CPU and GPU.

This parallelized a lot of functions, but most importantly the matrix vector product. At first we had a $(n_i, n_j) = (60, 60)$ grid, there was no significant speedup. We then tried a larger grid of 200x200 where we timed 59s on CPU and 9.73s on the GPU. This is about a 6x speedup.

For larger grids CG did not converge very well, so we preconditioned it with a Jacobi Preconditioner as in section 2.2. The convergence was much faster now, so to accurately measure speed we simply solve the equation around 100 times. Our results for comparing the CPU and the PyTorch is in figure 2. For small grid sizes, the CPU dominates and then as the grid size increases the power of the GPU is clearly visible. Note also that the plot has a logarithmic scale, we had a speedup of 8.1x with Pytorch compared to CPU on the finest grid.

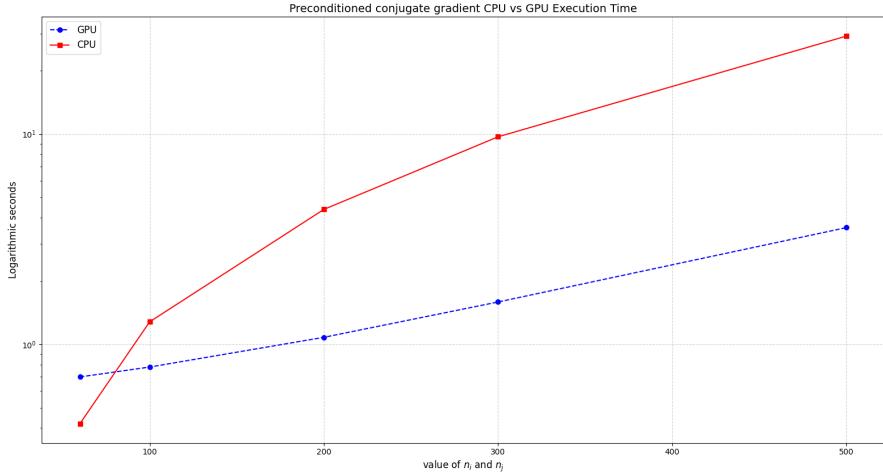


Figure 2: Computation time for cpu and gpu for various grid square grid sizes

A major flaw with PyTorch is that for the GPU version there is little support for sparse matrices, in particular at the time of writing, there is only support for Coordinate list (COO) matrix format on the GPU. Unfortunately, this format has a slow matrix vector product as it iterates through every element in the matrix for each element in the product. A format such as Compressed Sparse Row (CSR) format would efficiently compute the product with a better time complexity. Although the GPU beats the CPU there is still room for improvement.

5 Cupy Implementation

Cupy was easier to implement as Cupy is very similar to Numpy and Numpy is what the original code by Lada used. There is one major difference between the PyTorch and Cupy implementation and that is Cupy supports the Compressed Sparse Row format. When doing the Preconditioned CG algorithm the main bottleneck is the matrix dot product, but now we can efficiently compute it and we plot the PyTorch vs Cupy results in figure 3.

Cupy is performing better than PyTorch for every grid size that we tried. On the largest grid size of 1000×1000 we had a 9x speedup. Most importantly we can clearly see that Cupy has a better time complexity than the PyTorch implementation. This is due to the different sparse matrix formats discussed in an earlier section. This should be one polynomial degree time complexity difference as COO iterates every element while CSR iterates only the necessary elements.

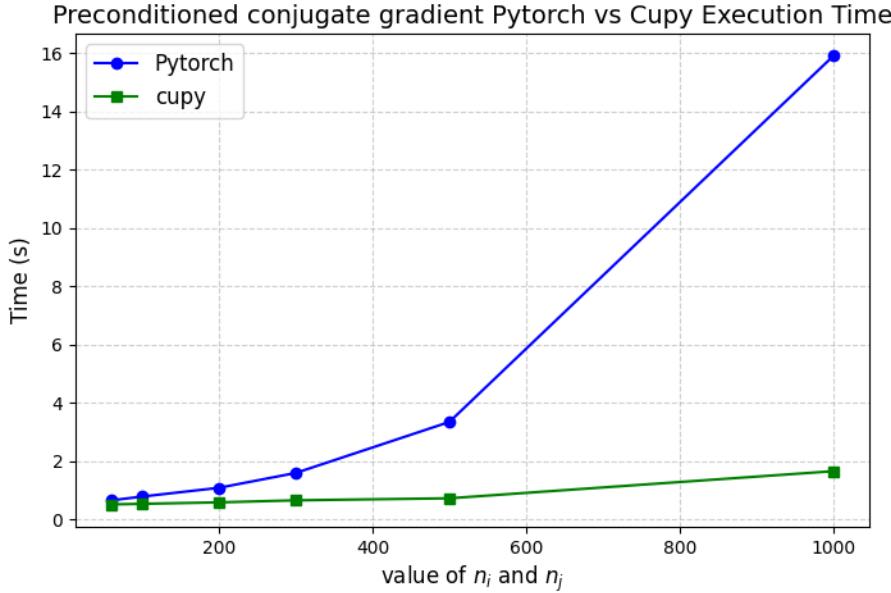


Figure 3: Computation time for PyTorch and Cupy for various grid square grid sizes

6 Multi-GPU Cupy

To achieve parallelization in the multi-GPU implementation, we employ domain decomposition. The computational domain is divided into subdomains, each assigned to a different GPU. This allows the solver to run concurrently on multiple GPUs, with each GPU handling a portion of the domain.

6.1 Subdomain and Ghost Cells

Decomposition is performed along the dimension by distributing the rows of the grid to available GPUs. For a domain of size $n_i \times n_j$, distributed across n GPUs, each GPU processes approximately $(n_i/n) \times n_j$ grid points. The subdomains are represented by instances of the GPUDomain class, which encapsulates the data and operations specific to each subdomain.

To ensure proper communication between subdomains, we introduce ghost cells [2]. Ghost cells are additional cells surrounding each subdomain that store a copy of the boundary data from neighboring subdomains⁴. These cells prompt the exchange of boundary information and maintain consistency across the global domain. When we initialize the subdomain of GPU, we need to allocate memory space for both subdomains and ghost cells. The number of ghost cells depends on the size of the finite difference stencil.

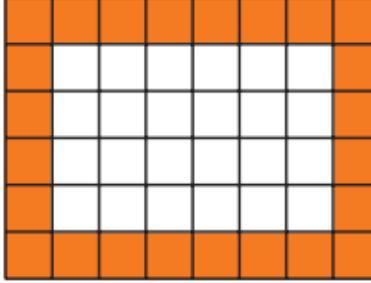


Figure 4: The subdomain(white part) and the ghost cell(orange part)

In the program, we use a 5 point stencil [3], which requires one layer of ghost cell on some side of the subdomain.

The Poisson equation in 2D can be discretized using the following formula:

$$u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1} - 4u_{i,j} = 0$$

This formula represents the relationship between a central point $u_{i,j}$ and its four neighboring points $(u_{i+1,j})$, $(u_{i-1,j})$, $(u_{i,j+1})$, and $(u_{i,j-1})$ in a uniform 2D grid. It need ghost cells to handle boundary conditions and ensure the accuracy of finite difference formulas near boundaries. Ghost cells are also an important tool for implementing parallelization.

6.2 Exchange boundary data

After every iteration of the solver, the boundary data of each neighboring subdomain need to be exchanged to update the ghost cells. All neighboring subdomain pairs should be traversed and the boundary data from one subdomain should be copied to the ghost cell of its neighbor subdomain, which ensures all the ghost cell have been synchronized with the latest data.

To optimize the inter-GPU communication, we use CuPy’s asynchronous copy functions (`cupy.cuda.runtime.memcpy_async`). These functions enable efficient data transfer between different GPU devices, minimizing the overhead of boundary exchange.

In the solver ghost cells are treated as part of the subdomain. But we need to treat the edges of the subdomain that lie on the edge of the grid differently. The ghost cells of those subdomains are expanding beyond the original grid. When we need to exchange boundary data, these ghost cells expanding beyond the grid are updated according to their own boundary conditions.

By managing the ghost cells and boundary exchange process, the multi-GPU implementation ensures seamless communication between subdomains and the consistency of the solution across the whole grid.

6.3 Result and Conclusion

The performance impact of the multi-GPU implementation is significant, especially for large problem sizes. For the case with 1000×1000 grid and maximum iteration count of 500 (maxit = 500), the version with one A40 GPU uses about 10.75s while the version with 2 A40 GPUs use about 6s.

The multi-GPU implementation shows the scalability of the Poisson solver and the ability to leverage the power of multiple GPUs to solve computationally intensive problems. By efficiently using domain decomposition, ghost cells, and optimized inter-GPU communication, we can significantly reduce the overall runtime and enable the solver to handle larger problem sizes in a fraction of the time compared to a single-GPU approach.

7 Problems encountered

Initially after converting to PyTorch there was a problem that went unnoticed for a while. The code was running and we were able to benchmark GPU versus CPU. However, the residual term

$$\epsilon = \mathbf{b} - \mathbf{Ax}_k$$

converged badly and plateaued after about 800 CG iterations with a big residue. The fact that the residue term did not decrease led us thinking the conjugate gradient method was not good enough. Therefore we attempted to solve this by trying other solvers and still it would not converge. It was not until a meeting with Lada, that we realized the issue was within the initialization of the stiffness matrix \mathbf{A} and not the actual solver. After fixing this issue, the solution converged as expected.

However, when the grid size was large such as 1000×1000 , CG was very unstable and slow to converge. The solution to this was to implement a preconditioner. We tried a Jacobi preconditioner and it solved the stability issue.

8 Implementation Improvements

Several key improvements were implemented to optimize the GPU-accelerated solver.

8.1 Sparse Matrix Operations

We implemented efficient sparse matrix construction and operations using PyTorch's sparse tensor capabilities. There are some key improvements.

- Efficient creation of sparse matrices using COO format
- Optimized matrix-vector multiplication using `torch.sparse.mm`

- Masked operations for boundary conditions to avoid unnecessary computations

8.2 Memory Management

Directly load data to GPU help us to minimize the data transfer between CPU and GPU.

```
x2d = torch.zeros((ni+1,nj+1), device=device)
```

This code also ensure the efficient use of some in-place to reduce memory allocation.

References

- [1] L. Davidson, “pypoisson,” 2024. [Online]. Available: <https://www.tfd.chalmers.se/~lada/pyPoisson.html>
- [2] Z. H. Ma, L. Qian, D. M. Causon, H. B. Gu, and C. G. Mingham, “A cartesian ghost-cell multigrid poisson solver for incompressible flows,” *International Journal for Numerical Methods in Engineering*, vol. 85, no. 2, pp. 230–246, 2011. [Online]. Available: <https://onlinelibrary.wiley.com/doi/abs/10.1002/nme.2967>
- [3] Wikipedia contributors, “Five-point stencil — Wikipedia, the free encyclopedia,” 2025, [Online; accessed 31-January-2025]. [Online]. Available: https://en.wikipedia.org/wiki/Five-point_stencil

Code

```

1 import torch
2 import numpy as np
3 import time
4 import socket
5 from scipy.sparse import spdiags, eye
6 import line_profiler as lp
7
8 profile_on = False
9 gpu = True
10
11 if profile_on:
12     profile = lp.LineProfiler()
13 else:
14     # If we are not profiling then define profile to
15     # ↪ just passthrough
16     def profile(func):

```

```

16         return func
17
18 if gpu and torch.cuda.is_available():
19     device = torch.device('cuda')
20 elif not gpu:
21     device = torch.device('cpu')
22 else:
23     print("Cuda is not available!")
24     exit()
25
26 print(f"Using device: {device}")
27
28
29 def setup_case():
30     global convergence_limit_u, dist, fx, fy, imon,
31             ↪ jmon, maxit, \
32     ni, nj, nsweep_u, solver_u, sormax, u_bc_east,
33             ↪ u_bc_east_type, u_bc_north, u_bc_north_type,
34             ↪ u_bc_south, u_bc_south_type, u_bc_west, \
35     u_bc_west_type, urf_u, viscos, vol, x2d, xp2d, y2d
36             ↪ , yp2d, device
37
38     ##### section 4 fluid properties #####
39     viscos = 1/10
40
41     ##### section 5 relaxation factors
42     ↪ #####
43     urf_u = 0.5
44
45     ##### section 6 number of iteration and
46             ↪ convergence criterira #####
47     maxit = 100
48     sormax = 1e-20
49
50     solver_u = 'cg' # Using PyTorch's conjugate
51             ↪ gradient solver
52     nsweep_u = 50
53     convergence_limit_u = 1e-6
54
55     ##### section 7 monitoring point #####
56     imon = ni-10
57     jmon = int(nj/2)
58
59     ##### section 10 boundary conditions
60             ↪ #####
61     # Boundary conditions for u (converted to tensors)

```

```

54     u_bc_west = torch.zeros(nj, device=device)
55     u_bc_east = torch.zeros(nj, device=device)
56     u_bc_south = torch.zeros(ni, device=device)
57     u_bc_north = torch.ones(ni, device=device) # =1
58         ↪ at north boundary, y=1
59
60     u_bc_west_type = 'd'
61     u_bc_east_type = 'd'
62     u_bc_south_type = 'd'
63     u_bc_north_type = 'd'
64
65     return
66
67 def init():
68     print('hostname:', socket.gethostname())
69
70     global x2d, y2d, xp2d, yp2d, dist, fx, fy, vol,
71         ↪ areaw, areas
72
73     # Distance to nearest wall
74     ywall_s = 0.5*(y2d[0:-1,0] + y2d[1:,0])
75     dist_s = yp2d - ywall_s.unsqueeze(1)
76     ywall_n = 0.5*(y2d[0:-1,-1] + y2d[1:,-1])
77     dist_n = ywall_n.unsqueeze(1) - yp2d
78
79     dist = torch.minimum(dist_s, dist_n)
80
81     # West face coordinate
82     xw = 0.5*(x2d[0:-1,0:-1] + x2d[0:-1,1:])
83     yw = 0.5*(y2d[0:-1,0:-1] + y2d[0:-1,1:])
84
85     del1x = ((xw-xp2d)**2 + (yw-yp2d)**2)**0.5
86
87     del2x = ((xw-torch.roll(xp2d, 1, dims=0))**2 +
88               (yw-torch.roll(yp2d, 1, dims=0))**2)**0.5
89     fx = del2x/(del1x + del2x)
90
91     # South face coordinate
92     xs = 0.5*(x2d[0:-1,0:-1] + x2d[1:,0:-1])
93     ys = 0.5*(y2d[0:-1,0:-1] + y2d[1:,0:-1])
94
95     del1y = ((xs-xp2d)**2 + (ys-yp2d)**2)**0.5
96     del2y = ((xs-torch.roll(xp2d, 1, dims=1))**2 +
97               (ys-torch.roll(yp2d, 1, dims=1))**2)**0.5
98     fy = del2y/(del1y + del2y)

```

```

98     # Area calculations
99     areawy = torch.diff(x2d, dim=1)
100    areawx = -torch.diff(y2d, dim=1)
101    areasy = -torch.diff(x2d, dim=0)
102    areasx = torch.diff(y2d, dim=0)
103
104    areaw = torch.sqrt(areawx**2 + areawy**2)
105    areas = torch.sqrt(areasx**2 + areasy**2)
106
107    # Volume calculation
108    ax = torch.diff(x2d, dim=1)
109    ay = torch.diff(y2d, dim=1)
110    bx = torch.diff(x2d, dim=0)
111    by = torch.diff(y2d, dim=0)
112
113    areaz_1 = 0.5*torch.abs(ax[0:-1,:,:]*by[:,0:-1] - ay
114        ↪ [0:-1,:,:]*bx[:,0:-1])
114    areaz_2 = 0.5*torch.abs(ax[1:,:,:]*by[:,1:] - ay
115        ↪ [1:,:,:]*bx[:,1:])
116    vol = areaz_1 + areaz_2
116
117    # Boundary coefficients
118    as_bound = areas[:,0]**2/(0.5*vol[:,0])
119    an_bound = areas[:, -1]**2/(0.5*vol[:, -1])
120    aw_bound = areaw[0,:,:]**2/(0.5*vol[0,:])
121    ae_bound = areaw[-1,:,:]**2/(0.5*vol[-1,:])
122
123    return areaw, areawx, areawy, areas, areasx,
124        ↪ areasy, vol, fx, fy, aw_bound, ae_bound,
125        ↪ as_bound, an_bound, dist
124
125 @profile
126 def solve_2d(phi2d, aw2d, ae2d, as2d, an2d, su2d, ap2d
127     ↪ , tol_conv, nmax):
127     # Convert inputs to PyTorch tensors
128     phi = torch.flatten(phi2d)
129     aw = torch.flatten(aw2d)
130     ae = torch.flatten(ae2d)
131     as1 = torch.flatten(as2d)
132     an = torch.flatten(an2d)
133     ap = torch.flatten(ap2d)
134     su = torch.flatten(su2d)
135
136     # Build sparse tensor
137     n = ni * nj
138     indices = []

```

```

139     values = []
140
141     # Main diagonal
142     diag_idx = torch.arange(n, device=device)
143     indices.append(torch.stack([diag_idx, diag_idx]))
144     values.append(ap)
145
146     # West coefficient (aw)
147     i = torch.arange(1, n, device=device)
148     j = i - 1
149     mask = (i % nj != 0)
150     i, j = i[mask], j[mask]
151     indices.append(torch.stack([i, j]))
152     values.append(aw[1:][mask])
153
154     # East coefficient (ae)
155     i = torch.arange(0, n-1, device=device)
156     j = i + 1
157     mask = ((i+1) % nj != 0)
158     i, j = i[mask], j[mask]
159     indices.append(torch.stack([i, j]))
160     values.append(ae[:-1][mask])
161
162     # South coefficient (as)
163     i = torch.arange(nj, n, device=device)
164     j = i - nj
165     indices.append(torch.stack([i, j]))
166     values.append(as1[nj:])
167
168     # North coefficient (an)
169     i = torch.arange(0, n-nj, device=device)
170     j = i + nj
171     indices.append(torch.stack([i, j]))
172     values.append(an[:-nj])
173
174     # Concatenate all indices and values
175     indices = torch.cat(indices, dim=1)
176     values = torch.cat(values)
177
178     # Create sparse tensor
179     A = torch.sparse_coo_tensor(indices, values, (n,n)
180                                ↪ , device=device)
181
182     # Define matrix-vector product function for
183     ↪ conjugate gradient
184     @profile

```

```

183     def mv(v):
184         return torch.sparse.mm(A, v.unsqueeze(1)).
185             → squeeze(1)
186
187     # Conjugate gradient solver
188     @profile
189     def conjugate_gradient(A_mv, b, x0, tol=1e-10,
190         ↪ max_iter=1000):
191         x = x0
192         r = b - A_mv(x)
193         p = r
194         r_norm_sq = r @ r
195
196         for _ in range(max_iter):
197             Ap = A_mv(p)
198             alpha = r_norm_sq / (p @ Ap)
199             x = x + alpha * p
200             r_new = r - alpha * Ap
201             r_norm_sq_new = r_new @ r_new
202             beta = r_norm_sq_new / r_norm_sq
203             r = r_new
204             r_norm_sq = r_norm_sq_new
205
206             if r_norm_sq < tol:
207                 break
208
209         p = r + beta * p
210
211     return x
212
213     # Initial guess
214     x0 = torch.zeros_like(su)
215
216     # Solve using conjugate gradient
217     phi = conjugate_gradient(mv, su, x0, tol=tol_conv,
218         ↪ max_iter=nmax)
219
220     # Reshape solution back to 2D
221     phi2d = phi.reshape(ni, nj)
222
223     # Calculate residual
224     resid = torch.norm(mv(phi) - su)
225
226     return phi2d, resid

```

```

226 def save_data(u2d):
227     print('save_data called')
228     # Convert tensor to numpy before saving
229     u_numpy = u2d.cpu().numpy()
230     np.save('u2d_saved', u_numpy)
231     return
232
233 # Main execution
234 if __name__ == "__main__":
235     # Load grid data and convert to tensors
236     datax = torch.from_numpy(np.loadtxt("x2d.dat")).to(
237         ↪ (device))
238     x = datax[0:-1]
239     ni = int(datax[-1].item())
240
241     datay = torch.from_numpy(np.loadtxt("y2d.dat")).to(
242         ↪ (device))
243     y = datay[0:-1]
244     nj = int(datay[-1].item())
245
246     # Initialize tensors on GPU
247     x2d = torch.zeros((ni+1,nj+1), device=device)
248     y2d = torch.zeros((ni+1,nj+1), device=device)
249
250     x2d = x.reshape(ni+1,nj+1)
251     y2d = y.reshape(ni+1,nj+1)
252
253     # Compute cell centers
254     xp2d = 0.25*(x2d[0:-1,0:-1] + x2d[0:-1,1:] + x2d
255                   ↪ [1:,0:-1] + x2d[1:,1:])
256     yp2d = 0.25*(y2d[0:-1,0:-1] + y2d[0:-1,1:] + y2d
257                   ↪ [1:,0:-1] + y2d[1:,1:])
258
259     # Initialize solution variables as tensors
260     u2d = torch.ones((ni,nj), device=device) * 1e-20
261
262     setup_case()
263     areaw, areawx, areawy, areas, areasx, areasy, vol,
264         ↪ fx, fy, aw_bound, ae_bound, as_bound,
265         ↪ an_bound, dist = init()
266
267     @profile
268     def coeff():
269         """GPU version of coefficient calculation"""
270         # Initialize viscosity tensors on GPU

```

```

265     visw = torch.ones((ni+1,nj), device=device) *
266         ↪ viscos
267     viis = torch.ones((ni,nj+1), device=device) *
268         ↪ viscos
269
270     # Initialize volume tensors
271     volw = torch.ones((ni+1,nj), device=device) * 1e
272         ↪ -10
273     vols = torch.ones((ni,nj+1), device=device) * 1e
274         ↪ -10
275
276     # Calculate volumes and diffusion coefficients
277     volw[1:,:] = 0.5 * torch.roll(vol, -1, dims=0) +
278         ↪ 0.5 * vol
279     diffw = visw[0:-1,:] * areaw[0:-1,:]**2 / volw
280         ↪ [0:-1,:]
281
282     vols[:,1:] = 0.5 * torch.roll(vol, -1, dims=1) +
283         ↪ 0.5 * vol
284     diffs = viis[:,0:-1] * areas[:,0:-1]**2 / vols
285         ↪ [:,0:-1]
286
287     # Set coefficients
288     aw2d = diffw
289     ae2d = torch.roll(diffw, -1, dims=0)
290     as2d = diffs
291     an2d = torch.roll(diffs, -1, dims=1)
292
293     # Zero out boundary coefficients
294     as2d[:,0] = 0
295     an2d[:, -1] = 0
296
297     # Initialize source terms
298     su2d = torch.zeros((ni,nj), device=device)
299     sp2d = torch.zeros((ni,nj), device=device)
300
301     if iter == 0:
302         print('aw[5,5],ae,as,an', aw2d[5,5].item(),
303               ↪ ae2d[5,5].item(),
304               ↪ as2d[5,5].item(), an2d[5,5].item())
305
306     return aw2d, ae2d, as2d, an2d, su2d, sp2d
307
308 @profile
309 def calcu(su2d, sp2d, aw2d, ae2d, as2d, an2d):

```

```

301     """GPU version of u-momentum source terms
302     ↪ calculation"""
303
304     # Add sources and modify source terms
305     su2d, sp2d = modify_u(su2d, sp2d)
306
307     # Calculate ap coefficient
308     ap2d = aw2d + ae2d + as2d + an2d - sp2d
309
310     # Apply under-relaxation
311     ap2d = ap2d / urf_u
312     su2d = su2d + (1-urf_u) * ap2d * u2d
313
314     return su2d, sp2d, ap2d
315
316 @profile
317 def bc(su2d,sp2d,phi_bc_west,phi_bc_east,phi_bc_south,
318       ↪ phi_bc_north\,
319       ,phi_bc_west_type,phi_bc_east_type,
320       ↪ phi_bc_south_type,phi_bc_north_type):
321
322     su2d = torch.zeros((ni,nj), device=device)
323     sp2d = torch.zeros((ni,nj), device=device)
324
325     #south
326     if phi_bc_south_type == 'd':
327         sp2d[:,0] = sp2d[:,0] - viscos*as_bound
328         su2d[:,0] = su2d[:,0] + viscos*as_bound*
329             ↪ phi_bc_south
330
331     #north
332     if phi_bc_north_type == 'd':
333         sp2d[:, -1] = sp2d[:, -1] - viscos*an_bound
334         su2d[:, -1] = su2d[:, -1] + viscos*an_bound*
335             ↪ phi_bc_north
336
337     #west
338     if phi_bc_west_type == 'd':
339         sp2d[0,:] = sp2d[0,:] - viscos*aw_bound
400         su2d[0,:] = su2d[0,:] + viscos*aw_bound*
401             ↪ phi_bc_west
402
403     #east
404     if phi_bc_east_type == 'd':
405         sp2d[-1,:] = sp2d[-1,:] - viscos*ae_bound

```

```

340         su2d[-1,:] = su2d[-1,:] + viscos*ae_bound*
341             ↪ phi_bc_east
342
343     return su2d, sp2d
344
345 @profile
346 def modify_u(su2d, sp2d):
347     # Add a point source/volume of 100 at x = 1.5 and
348     ↪ y = 0.5
349     xx = 1.5
350
351     i1 = torch.argmin(torch.abs(xx-xp2d[:,1]))
352     yy = 0.5
353     j1 = torch.argmin(torch.abs(yy-yp2d[1,:]))
354     su2d[i1,j1] = su2d[i1,j1] + 100*vol[i1,j1]
355
356     return su2d, sp2d
357
358 @profile
359 def main_loop(u2d):
360     cumulative_time = 0
361     # Main iteration loop
362     for iter in range(maxit):
363         start_time_iter = time.time()
364         start_time = time.time()
365         # Compute coefficient matrices
366         aw2d, ae2d, as2d, an2d, su2d, sp2d = coeff()
367
368         # Apply boundary conditions for u2d
369         su2d, sp2d = bc(su2d, sp2d, u_bc_west,
370                         ↪ u_bc_east, u_bc_south, u_bc_north,
371                         ↪ u_bc_west_type, u_bc_east_type
372                         ↪ , u_bc_south_type,
373                         ↪ u_bc_north_type)
374
375         # Calculate source terms for u momentum
376         ↪ equation
377         su2d, sp2d, ap2d = calcu(su2d, sp2d, aw2d,
378             ↪ ae2d, as2d, an2d)
379
380         # Solve for u using conjugate gradient method
381         if maxit > 0:
382             u2d, residual_u = solve_2d(u2d, aw2d, ae2d
383                 ↪ , as2d, an2d, su2d, ap2d,
384                 ↪ convergence_limit_u
385                 ↪ , nsweep_u)

```

```

377
378     # Print computation time for u
379     print(f"{'time_u:'}{time.time()-start_time:.2
380         ↪ e}")
380     cumulative_time += time.time()-start_time
381
382     # Monitor convergence
383     if iter % 10 == 0:
384         print(f"\n--iter:{iter:d}, residual:{{
385             ↪ residual_u:.2e}\n")
385         print(f"\nmonitor_iteration:{iter:4d}, u:{{
386             ↪ u2d[imon,jmon].item():.2e}\n")
387
388     # Calculate maximum velocity for stability
389     # monitoring
390     umax = torch.max(u2d).item()
391     print(f"\n---iter:{iter:2d}, umax:{umax:.2e
392         ↪ }\n")
393
394     # Print iteration timing
395     print(f"time_one_iteration:{time.time()-
396         ↪ start_time_iter:.2e}")
397
398     # Check convergence
399     if residual_u < sormax:
400         break
401
402     return cumulative_time
403
404 total_time = main_loop(u2d)
405
406 # Print and save profiler results
407 if profile_on:
408     profile.print_stats()
409     with open('profile_results.txt', 'w') as f:
410         profile.print_stats(stream=f)
411
412 # Save final results
413 save_data(u2d)
414 print('program_reached_normal_stop')
415 print(f"{'Cumulative_time:'}{total_time}")

```